

# Numerical Optimisation. Project 1

## Team Information

*Group 1*

*Participants information in alphabetical order*

#	Name	Lastname	Matr Number
1	David	Körnsteiner	11820336
2	Christian	Peinthor	11815592
3	Elias	Ramoser	11918558
4	Georg	Storz	11918811

## Implementation

*All points  $x$  are represented as numpy arrays. Function  $f$  returns a scalar with  $\text{grad}_f$  and  $\text{hessian}_f$  returning numpy arrays.*

---

## Imports

*Describe how to install additional packages, if you have some, here*

```
In [1]: import numpy as np
import scipy
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
from numpy.linalg import *
from sklearn.datasets import make_spd_matrix
```

---

## Stopping criteria

*Function returns True if the gradient of  $f$  at  $x_k$  relative to  $x_0$  is smaller than parameter  $\text{tol}$ .*

Additionally there is an upper bound for iterations to stop non converging algorithms.

In [2]:

```
def stop_crit(grad_f, xk, x0, i, tol=1e-8, max_iter=5000):  
    if i > max_iter:  
        return True  
    elif norm(grad_f(xk)) <= tol * norm(grad_f(x0)):  
        return True  
    return False
```

---

## Varibales scaling

Place your reasoning here, how your algorithm behave with respect to this problem. You can also try rescaling your problems This is additional task, which can earn you several points.

In [ ]:

---

## Stabilising algorithm

Place your reasoning here, how your algorithm behave with respect to this problem. You can also try rescaling your problems This is additional task, which can earn you several points.

In [3]:

```
#your function for stabilising goes here
```

---

## Fighting floating-point numbers and roundoff error

Place your reasoning, how your algorithm behave with respect to this problem. You can also try rescaling your problems This is additional task, which can earn you several points.

---

## Inverting matrices

`linear_solve()` provides a way to solve linear systems of equations using a LU-factorization of  $A$  and subsequent forward and backward substitution as described in the book. This solver proves to be quite unstable though in practical applications. We therefore use the numpy implementation of `solve()`.

In [4]:

```
def forward_substitution(L, b):  
    n = L.shape[0]
```

```

y = np.zeros_like(b, dtype=np.double);

y[0] = b[0] / L[0, 0]

for i in range(1, n):
    y[i] = (b[i] - np.dot(L[i,:i], y[:i])) / L[i,i]

return y

def back_substitution(U, y):

    n = U.shape[0]

    x = np.zeros_like(y, dtype=np.double);

    x[-1] = y[-1] / U[-1, -1]

    for i in range(n-2, -1, -1):
        x[i] = (y[i] - np.dot(U[i,i:], x[i:])) / U[i,i]

    return x

def linear_solve(A, b):

    P, L, U = scipy.linalg.lu(A)

    y = forward_substitution(L, P @ b)

    return back_substitution(U, y)

```

## Gradients calculation

Following functions are wrapper functions that provide approximations of the gradient and hessian of  $f$  using the forward-difference approach as described in the book.

In [5]:

```

def e_i(size, index):
    arr = np.zeros(size)
    arr[index] = 1.0
    return arr

def approx_grad(f, e=1.1*10**-8):
    def grad_f(x):
        if x.size == 1:
            return (f(x + e) - f(x)) / e
        return np.array([(f(x + e * e_i(x.size, i)) - f(x)) / e for i in range(x.size)])
    return grad_f

def approx_hessian(f, e=1.1*10**-8):
    def hessian_f(x):
        if x.size == 1:
            return (f(x + 2*e) - 2*f(x + e) + f(x)) / e**2
        return np.array([(f(x + e * e_i(x.size, i) + e * e_i(x.size, j)) - f(
            x + e * e_i(x.size, i)) - f(x + e * e_i(x.size, j)) + f(
            x)) / e**2 for j in range(x.size)] for i in range(x.size)])
    return hessian_f

```

## Additional objects you implemented

The class `Problem()` provides an object to generate and set up quadratic and non quadratic test problems for the algorithms with.

In [6]:

```
class Problem():

    def __init__(self):

        self.f = None
        self.grad_f = None
        self.hessian_f = None
        self.min_x = None

    def quadratic(self, n_dim, rseed):

        rng = np.random.RandomState(rseed)
        A = make_spd_matrix(n_dim, random_state=rseed)
        x = rng.randint(-10, 10, n_dim)
        b = A @ x

        def f(x):
            return 0.5 * x.T @ A @ x - b @ x

        def grad_f(x):
            return A @ x - b

        def hessian_f(x):
            return A

        self.f = f
        self.grad_f = grad_f
        self.hessian_f = hessian_f
        self.min_x = x
        self.A = A
        self.b = b

    def rosenbrock(self):

        def f(x):
            return 100*(x[1] - x[0]**2)**2 + (1 - x[0])**2

        def grad_f(x):
            return np.array([-400*x[0]*(x[1] - x[0]**2) - 2*(1 - x[0]),
                             200*(x[1] - x[0]**2)])

        def hessian_f(x):
            return np.array([[ -400*(x[1] - 3*x[0]**2) + 2, -400*x[0]],
                             [-400*x[0], 200]])

        self.f = f
        self.grad_f = grad_f
        self.hessian_f = hessian_f
        self.min_x = np.array([1,1])

    def himmelblau(self):

        def f(x):
            return (x[0]**2 + x[1] - 11)**2 + (x[0] + x[1]**2 - 7)**2

        def grad_f(x):
```

```

        return np.array([4*x[0]*(x[0]**2 + x[1] - 11)+2*(x[0] + x[1]**2 - 7),
                        4*x[1]*(x[1]**2 + x[0] - 7)+2*(x[1] + x[0]**2 - 11)])

def hessian_f(x):
    return np.array([[12*x[0]**2 + 4*x[1] - 42, 4*(x[1] + x[0])],
                    [4*(x[1] + x[0]), 12*x[1]**2 + 4*x[0] - 26]])

self.f = f
self.grad_f = grad_f
self.hessian_f = hessian_f
self.min_x = np.array([[3,2], [-2.805118, 3.131312], [-3.779310, -3.283186], [3.

def poly_1(self):

    def f(x):
        return ((x - 7)**2 * (x - 3)**2) / 4

    def grad_f(x):
        return (x - 7) * (x - 5) * (x - 3)

    def hessian_f(x):
        return 3 * x**2 - 30 * x + 71

    self.f = f
    self.grad_f = grad_f
    self.hessian_f = hessian_f
    self.min_x = np.array([[3],[5],[7]])

def poly_2(self):

    def f(x):
        return (x**2 * (x**2 - 16*x + 40)) / 4

    def grad_f(x):
        return x * (x - 2) * (x - 10)

    def hessian_f(x):
        return 3 * x**2 - 24 * x + 20

    self.f = f
    self.grad_f = grad_f
    self.hessian_f = hessian_f
    self.min_x = np.array([[0],[2],[10]])

def poly_3(self):

    def f(x):
        return (x * (3 * x**3 - 64 * x**2 + 414 * x - 648)) / 12

    def grad_f(x):
        return (x - 1) * (x - 6) * (x - 9)

    def hessian_f(x):
        return 3 * x**2 - 32 * x + 69

    self.f = f
    self.grad_f = grad_f
    self.hessian_f = hessian_f
    self.min_x = np.array([[1],[6],[9]])

```

`alpha_wolfe()` returns a step length satisfying the weak wolfe conditions using a bisection approach as described in [1]. `SR1()` implements the line search quasi newton method utilizing SR1 updating for inverse Hessian approximation and a method for stabilizing that resets  $H$  as a multiple of  $I$  inspired by the method deployed in [2].

In [7]:

```
def alpha_wolfe(f, grad_f, xk, pk, c1=1e-4, c2=0.9):

    alpha = 0
    beta = np.Inf
    t = 1

    while True:

        if f(xk + t * pk) > (f(xk) + c1 * t * (pk @ grad_f(xk))):
            beta = t
            t = 0.5 * (alpha + beta)
        elif (-pk @ grad_f(xk + t * pk)) > (-c2 * pk @ grad_f(xk)):
            alpha = t
            t = (2 * alpha if beta == np.Inf else 0.5 * (alpha + beta))
        else:
            return t

def SR1(x0 : np.array, f, grad_f=None, r=1e-8):

    conv_tol = 1e-8
    if grad_f == None:
        grad_f = approx_grad(f)
        conv_tol = 1e-6

    i = 0
    H = np.identity(x0.size)
    x = x0

    while not stop_crit(grad_f, x, x0, i, tol=conv_tol):

        pk = - H @ grad_f(x)
        x_1 = x + alpha_wolfe(f, grad_f, x, pk) * pk
        sk = x_1 - x
        yk = grad_f(x_1) - grad_f(x)
        rhok = 1 / (yk.T @ sk)

        if ((sk @ yk - yk @ H @ yk) < 0) or (abs(yk @ (sk - H @ yk)) < r * np.linalg.norm(H, np.inf) > 1e10):
            mu = (sk @ sk) / (yk @ sk) - ((sk @ sk)**2 / (yk @ sk)**2 - (sk @ sk) / (yk @ H_1 @ sk)) / (yk @ sk)
            H_1 = mu * np.identity(x0.size)
            H = H_1
        else:
            H_1 = H + np.outer((sk - H @ yk), (sk - H @ yk)) / ((sk - H @ yk) @ yk)
            H = H_1
        x = x_1
        i += 1

    print(f"\nsearch terminated at iteration {i} | result: {x}")

    return x
```

Testing on 5-10 variables, Quadratic objective

---

## Implement functions to optimise over

*Place for additional comments and argumentation*

```
In [8]: rseed = [1,4,6,7,8]
quadratic_probs = []
for i in range(5):
    prob = Problem()
    prob.quadratic(10, rseed[i])
    quadratic_probs.append(prob)
```

---

## Run 5 tests

**Note:** After every test print out the results.

For your convinience we implemented a function which will do it for you. Function can be used in case after running optimisation you return  $x_{optimal}$ , and if you have implemented your gradient approximation. Feel free to bring your adjustments.

Additionally print how many iterations your algorithm needed. You might also provide charts of your taste (if you want).

*Place for your additional comments and argumentation*

```
In [9]: def final_printout(x_0,x_optimal,x_appr,f,grad,args,tolerance):
        """
        Parameters
        -----
        x_0: numpy 1D array, corresponds to initial point
        x_optimal: numpy 1D array, corresponds to optimal point, which you know, or have
        x_appr: numpy 1D array, corresponds to approximated point, which your algorithm
        -----
        f: function which takes 2 inputs: x (initial, optimal, or approximated)
            **args
            Function f returns a scalar output.
        -----
        grad: function which takes 3 inputs: x (initial, optimal, or approximated),
            function f,
            args (which are submitted, because you migh
            to call f(x,**args) inside your gradie
            Function grad approximates gradient at given point and returns a 1d np arr
        -----
        args: dictionary, additional (except of x) arguments to function f
        tolerance: float number, absolute tolerance, precision to which, you compare opt
        """

        print(f'Initial x is :\t\t{x_0}')
        print(f'Optimal x is :\t\t{x_optimal}')
        print(f'Approximated x is :\t{x_appr}')
        print(f'Is close verificalion: \t{np.isclose(x_appr,x_optimal,atol=tolerance)}\n')
```

```

f_opt = f(x_optimal,**args)
f_appr = f(x_appr,**args)
print(f'Function value in optimal point:\t{f_opt}')
print(f'Function value in approximated point: \t{f_appr}')
print(f'Is close verificaion:\t{np.isclose(f_opt,f_appr,atol=tolerance)}\n')
print(f'Gradient approximation in optimal point is:\n{grad(f,x_optimal,args)}\n')
grad_appr = grad(f,x_appr,args)
print(f'Gradient approximation in approximated point is:\n{grad_appr}\n')
print(f'Is close verificaion:\n{np.isclose(grad_appr,np.zeros(grad_appr.shape),a

```

In [10]:

```

for i, prob in enumerate(quadratic_probs):
    print(f"Problem {i+1}: ")
    print("approximated gradient: ")
    SR1(np.zeros(10), prob.f)
    print("\n exact gradient: ")
    SR1(np.zeros(10), prob.f, prob.grad_f)
    print(f"\n actual minimum: {prob.min_x}\n")

```

Problem 1:

approximated gradient:

```

search terminated at iteration 58 | result: [-5.0001946   0.99989553  2.00008187 -1.
99983473 -1.00006189  1.00024194
-5.00007264  4.99971227 -9.99995911  5.99988959]

```

exact gradient:

```

search terminated at iteration 25 | result: [ -4.9999994   1.00000014  1.99999957
-2.00000079  -0.99999977
 0.99999874  -4.99999975  5.00000096 -10.00000026  6.00000034]

```

actual minimum: [ -5 1 2 -2 -1 1 -5 5 -10 6]

Problem 2:

approximated gradient:

```

search terminated at iteration 26 | result: [ 4.00024531 -5.00008354 -8.99991686 -2.
00011483 -2.00004278  7.99995868
-0.99993076 -2.99977594  3.00002912 -1.9998709 ]

```

exact gradient:

```

search terminated at iteration 14 | result: [ 4. -5. -9. -2. -2.  8. -1. -3.  3. -
2.]

```

actual minimum: [ 4 -5 -9 -2 -2 8 -1 -3 3 -2]

Problem 3:

approximated gradient:

```

search terminated at iteration 63 | result: [ 8.88034994e-05 -1.00010929e+00 -6.9999
8888e+00 -4.56988777e-06
 3.00010501e+00  4.99996078e+00  5.46075263e-05  5.99998395e+00
-8.99992978e+00  9.99913716e-01]

```

exact gradient:

```

search terminated at iteration 14 | result: [-8.29467571e-13 -1.00000000e+00 -7.0000
0000e+00 -4.20259313e-13
 3.00000000e+00  5.00000000e+00 -1.26430810e-12  6.00000000e+00
-9.00000000e+00  1.00000000e+00]

```

actual minimum: [ 0 -1 -7 0 3 5 0 6 -9 1]

Problem 4:

approximated gradient:



```
search terminated at iteration 96 | result: [ 5.00016106e+00 -5.99995237e+00 -6.9997
1208e+00  9.00007398e+00
-3.00001668e+00  3.99994430e+00 -1.99974390e+00  3.99990280e+00
 8.81652477e-05 -1.99990065e+00]
```

exact gradient:

```
search terminated at iteration 54 | result: [ 4.99999930e+00 -6.00000035e+00 -7.0000
0064e+00  8.99999980e+00
-3.00000024e+00  3.99999990e+00 -2.00000062e+00  3.99999994e+00
-5.86627627e-07 -2.00000047e+00]
```

actual minimum: [ 5 -6 -7 9 -3 4 -2 4 0 -2]

Problem 5:

approximated gradient:

```
search terminated at iteration 33 | result: [-7.00001463  6.99998531 -1.00006806 -5.
00014037 -2.00008892  9.00005268
-1.99999778  6.00006372  2.99993072  7.00002436]
```

exact gradient:

```
search terminated at iteration 14 | result: [-7.  7. -1. -5. -2.  9. -2.  6.  3.
 7.]
```

actual minimum: [-7 7 -1 -5 -2 9 -2 6 3 7]

The quasi newton method is the second best method, of the four under review. It can solve the problems in under 100 iteration in our test.

## Testing on functions of 1-2 variables, Non-quadratic objective

---

### Implement functions to optimise over

*Place for additional comments and argumentation*

In [ ]:

---

### Run 5 tests

*Place for your additional comments and argumentation*

In [11]:

```
prob = Problem()
prob.rosenbrock()
print(f"Problem rosenbrock: \n")
print("approximate gradient: ")
SR1(np.array([1.2,1.2]), prob.f)
```

```

print("\nexact gradient: ")
SR1(np.array([1.2,1.2]), prob.f, prob.grad_f)
print(f"\nactual minimum: {prob.min_x}")

prob = Problem()
prob.himmelblau()
print(f"\nProblem himmelblau: \n")
print("approximate gradient: ")
SR1(np.array([0,0]), prob.f)
print("\nexact gradient: ")
SR1(np.array([0,0]), prob.f, prob.grad_f)
print(f"\nactual minimum: {prob.min_x}")

prob = Problem()
prob.poly_1()
print(f"\nProblem poly_1: \n")
print("approximate gradient: ")
SR1(np.array([2]), prob.f)
print("\nexact gradient: ")
SR1(np.array([2]), prob.f, prob.grad_f)
print(f"\nactual minimum: {prob.min_x}")

prob = Problem()
prob.poly_2()
print(f"\nProblem poly_2: \n")
print("approximate gradient: ")
SR1(np.array([1]), prob.f)
print("\nexact gradient: ")
SR1(np.array([1]), prob.f, prob.grad_f)
print(f"\nactual minimum: {prob.min_x}")

prob = Problem()
prob.poly_3()
print(f"\nProblem poly_3: \n")
print("approximate gradient: ")
SR1(np.array([7]), prob.f)
print("\nexact gradient: ")
SR1(np.array([7]), prob.f, prob.grad_f)
print(f"\nactual minimum: {prob.min_x}")

```

Problem rosenbrock:

approximate gradient:

search terminated at iteration 22 | result: [0.9999907 0.99998138]

exact gradient:

search terminated at iteration 24 | result: [1. 1.]

actual minimum: [1 1]

Problem himmelblau:

approximate gradient:

search terminated at iteration 11 | result: [2.99999999 2. ]

exact gradient:

search terminated at iteration 11 | result: [3. 2.]

actual minimum: [[ 3. 2. ]  
 [-2.805118 3.131312]  
 [-3.77931 -3.283186]  
 [ 3.584428 -1.848126]]

*Problem poly\_1:*

*approximate gradient:*

*search terminated at iteration 7 | result: [6.99999993]*

*exact gradient:*

*search terminated at iteration 8 | result: [7.]*

*actual minimum: [[3]*

*[5]*

*[7]]*

*Problem poly\_2:*

*approximate gradient:*

*search terminated at iteration 5 | result: [-5.34496639e-09]*

*exact gradient:*

*search terminated at iteration 5 | result: [1.55022962e-10]*

*actual minimum: [[ 0]*

*[ 2]*

*[10]]*

*Problem poly\_3:*

*approximate gradient:*

*<ipython-input-7-da1c763f9de6>:39: RuntimeWarning: invalid value encountered in double\_scalars*

*mu = (sk @ sk) / (yk @ sk) - ((sk @ sk)\*\*2 / (yk @ sk)\*\*2 - (sk @ sk) / (yk @ yk))*

*\*\*0.5*

*search terminated at iteration 5001 | result: [nan]*

*exact gradient:*

*search terminated at iteration 9 | result: [9.]*

*actual minimum: [[1]*

*[6]*

*[9]]*

*The quasi newton method can calculate the minima of the non quadratic functions very fast. But it has problems with some of the functions*

## Template for teachers' tests

---

### Set up a template, how one can run your code

*Template should include skeletons for:*

- *custom function to optimise over*
- *values initialisation to submit into optimising algorithm*
- *optimiser function call*

- report print out call

Provide descriptions and comments.

In [12]:

```
algorithm_to_test = SR1

# Here you can set your individual starting point
x_0 = np.ones(10)

# Here you can enter your individual function
def f(x):
    return np.sum(np.square(x))

# Here you can enter the exact gradient of your function
# This function will just be used in the second test
def grad_f(x):
    return 2*x

# Test run:

print("Test with approximate gradient:")
algorithm_to_test(x_0, f)

print('\n'*2) # Print some lines between the tests

print("Test with exact gradient:")
algorithm_to_test(x_0, f, grad_f)

print()
```

Test with approximate gradient:

```
search terminated at iteration 1 | result: [-1.99687822e-09 -1.99687822e-09 -1.99687
822e-09 -1.99687822e-09
-1.99687822e-09 -1.99687822e-09 -1.99687822e-09 -1.99687822e-09
-1.99687822e-09 -1.99687822e-09]
```

Test with exact gradient:

```
search terminated at iteration 1 | result: [0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

```
<ipython-input-7-da1c763f9de6>:39: RuntimeWarning: invalid value encountered in doub
le_scalars
    mu = (sk @ sk) / (yk @ sk) - ((sk @ sk)**2 / (yk @ sk)**2 - (sk @ sk) / (yk @ yk))
    **0.5
```

## References

- [1] <https://sites.math.washington.edu/~burke/crs/408/notes/nlp/line.pdf>  
[2] <https://www.sciencedirect.com/science/article/pii/S0898122111004202?via%3Dihub>